

Avalux: An Open Integrated Platform for Moderated Usability Testing with Real-Time Evaluator–Participant Synchronization

Hugo Correia, Bernardo Marques, and Samuel Silva
IEETA / DETI, University of Aveiro
Aveiro, Portugal
{hf.correia, bernardo.marques, sss}@ua.pt

Abstract—Moderated usability testing remains the reference method for discovering interaction problems, yet its instrumentation is fragmented: evaluators juggle stopwatches, spreadsheets, questionnaire tools, and manual report assembly, while commercial platforms focus on unmoderated remote studies whose moderated modes capture video and free-form notes rather than structured per-task metrics. We present Avalux, an open, self-hostable web platform that integrates the complete moderated-testing pipeline: reusable test protocols (tasks with optimal-path baselines, typed error taxonomies, per-task questions including in-browser audio/video/photo capture, interview scripts, and study-specific participant attributes); a live evaluator cockpit that logs per-task completion time, action counts, categorized errors, hesitation events, and Single Ease Question ratings; and a participant client that runs on the participant’s own device via a join link, synchronized with the evaluator in real time over WebSocket channels, including a gating protocol that prevents the session from advancing while the participant is still answering. Sessions can conclude with a semi-structured interview and with whichever automatically scored questionnaires a study selects, among them the System Usability Scale, and export to a self-contained PDF report. For teaching, a further layer turns properties of the methods into figures students compute on their own work: independent heuristic inspection that reports how far evaluators agreed, a comparison of inspection predictions with what participants actually encountered, reflection after each session drawn from what the session recorded, a kept record of every protocol revision, and a class-wide view for the instructor. We describe the architecture, a React single-page application over a Postgres backend with row-level security, and report two case studies: a four-participant evaluation of a household touch-panel controller (mean SUS 28.1) and a two-participant evaluation of an eye-tracking analysis dashboard (SUS 80). The studies show that the platform carries complete, realistic evaluations end to end, including a physical device beyond the reach of browser-based instrumentation, while producing diagnostic task-level detail alongside the standard summary metrics.

Index Terms—usability testing, moderated evaluation, heuristic evaluation, SUS, SEQ, HCI education, research tools, human-computer interaction

I. INTRODUCTION

Task-based moderated testing with a small number of participants remains the workhorse of usability practice: a facilitator observes a participant attempting representative tasks, records performance and difficulties, and complements observation with standardized instruments such as the System Usability Scale (SUS) [1] and the Single Ease Question (SEQ) [2].

Classic problem-discovery models show that even four to five participants uncover the majority of severe problems [3], [4], which keeps the method attractive for small teams and academic projects.

The tooling situation is less attractive. Commercial platforms concentrate on unmoderated remote studies, where metrics are captured by automatic instrumentation of the participant’s browser; their moderated modes produce session videos and free-form notes that must be coded after the fact. In practice, a moderated session is therefore run with an ad-hoc toolchain (a stopwatch, a paper or spreadsheet observation grid, a separate questionnaire form, and manual report writing) that is slow, error-prone, and hard to standardize across studies.

This paper presents *Avalux*, an open web platform that integrates the moderated-testing pipeline end to end. Its contributions are:

- an **integrated protocol model**: reusable templates couple tasks (with optimal time and action-count baselines) to typed error taxonomies, per-task questions (open, choice, rating, and in-browser audio/video/photo capture), interview scripts, and study-specific participant attributes;
- a **live moderated session** design in which the evaluator’s cockpit (timer, action counter, one-tap typed error and hesitation logging, per-task SEQ) and the participant’s own device are synchronized in real time, including a gating protocol that blocks task advancement while the participant is still answering the previous task’s questions;
- **two case studies**, a household touch-panel controller and an eye-tracking analysis dashboard, both run entirely on the platform, showing that it carries realistic studies end to end, including a physical device beyond browser instrumentation, and that its task-level instruments localize the problems behind widely differing outcomes (mean SUS 28.1 vs. 80);
- a **pedagogical design** for classroom use, mapped onto empirically reported learning difficulties, in which features are chosen for the evidence they produce about the method rather than for the effort they save: independent heuristic inspection with an agreement statistic, a comparison of inspection predictions against observed problems, co-rating, evidence-grounded reflection, a protocol review and an optional instructor gate with a kept revision

history, and a class view for the instructor.

Avalux is deployed at <https://avalux.pt>; its schema, protocol seeds, and report generator are versioned with the application source.

II. RELATED WORK

A mature market of commercial software-as-a-service platforms now supports remote usability evaluation. UserTesting combines a proprietary participant panel with video-first moderated (“Live Conversation”) and unmoderated studies overlaid with AI-assisted analysis [5], and has absorbed the quantitative benchmarking platform UserZoom following their 2023 merger [6]. Maze is unmoderated by default, converting prototypes and live websites into task-based studies scored through automatically instrumented metrics such as completion rate, misclick rate, and task duration [7]. Lookback specializes in live moderated sessions with screen, camera, and audio capture, stakeholder observation rooms, and timestamped free-form notes [8]; Loop11 supports both modes and, unusually, auto-computes SUS and NPS for unmoderated tests [9]; Useberry targets prototype-centric unmoderated testing with click and flow analytics [10]. Across these tools, quantitative evidence derives almost entirely from automated client-side instrumentation in unmoderated studies, while moderated modes yield video recordings and free-form notes that must be coded after the fact. To the best of our knowledge, none *combines* structured real-time evaluator logging in moderated sessions, per-task ease ratings integrated into the same pipeline, and an open, self-hostable architecture in one platform.

Academic work on remote evaluation dates to Hartson et al., who framed the network as an extension of the usability laboratory [11]. Andreasen et al. subsequently found remote synchronous testing “virtually equivalent” to conventional lab-based think-aloud, whereas asynchronous methods identify fewer problems at lower moderator cost [12], an asymmetry confirmed by Bruun et al. [13] and persisting in recent comparisons [14]. On the tooling side, RemUSINE and WebRemUSINE pioneered task-level instrumentation by comparing interaction logs against task models, but operate post hoc rather than steering live sessions [15], [16]. TechSmith’s Morae became the de facto observation suite in usability practice [17], [18]; its Observer component supported real-time marker and error logging by observers during moderated sessions and is the closest precedent to our evaluator instrument, but Morae is a discontinued desktop product requiring installed software on the participant’s machine, without participant-side questioning or an open architecture. More recent browser-based platforms such as UTAssistant [19], [20] and the analytics-augmented test-management system of Generosi et al. [21] target largely asynchronous workflows, and an action-research deployment of remote testing found the evaluation tool itself to be an obstacle to high-quality results [22].

Avalux’s measurement model follows the ISO 9241-11 decomposition of usability into effectiveness, efficiency, and satisfaction [23]. Perceived usability is captured with the ten-item SUS [1], whose reliability [24], [25], benchmark mean

of 68 [26], and adjective-based interpretation [27] are well established; per-task difficulty uses the Single Ease Question, shown to perform comparably to more elaborate post-task instruments [2]. Completion rate, time-on-task, and action counts relative to an optimal path are the standard objective performance metrics [28], [29], and moderated testing with small samples is supported by classic problem-discovery models [3], [4].

A further thread concerns inspection and how evaluation methods are compared. Heuristic evaluation has evaluators judge an interface against a short set of principles without participants [30], and Nielsen’s set of ten heuristics [31] is the one Avalux builds in. Its results depend heavily on who applies it: across eleven studies of cognitive walkthrough, heuristic evaluation, and thinking-aloud, average agreement between any two evaluators ranged from 5% to 65% [32]. Comparisons between methods rest on the measures of thoroughness and validity defined by Sears [33] and systematized by Hartson et al., who show that any such comparison depends on the criterion chosen to decide which problems are real [34].

How HCI should be taught is a long-running discussion in the field [35]. A systematic mapping of HCI teaching found active, project-based approaches dominant, only two tools built specifically for teaching the subject, and a literature made up largely of experience reports rather than rigorous evaluations [36]; what students find difficult has since been documented empirically [37]. Tools for teaching evaluation support planning a study [38], [39] rather than running it, measuring the reliability of its findings, or comparing a method’s predictions with what testing then shows. Generative assistants, the most recent entrant, are valued by students as procedural scaffolds but distrusted when they pre-empt the students’ own judgement [40]. Contributions of this kind are commonly evaluated by demonstration, usage, technical performance, or heuristic analysis [41], and the Discussion places the present evaluation within that taxonomy.

A final thread concerns where sessions take place and on whose device. Early comparisons found no significant remote-versus-local differences in critical errors or time-on-task [42], [43], while lab-versus-field studies of mobile applications show that removing participants from their natural device context can mask context-dependent problems [44], [45]. The COVID-19 pandemic pushed moderated protocols onto general-purpose videoconferencing [46], [47], with studies of remote moderation identifying the moderator’s reduced visibility into the participant’s device and context as its principal cost [48].

Avalux addresses the gap these threads leave open. It is an open, integrated platform for moderated testing in which the evaluator live-logs, per task, completion time, action counts, typed errors, hesitation events, and SEQ ratings, while the participant completes tasks on their own device via a browser join link, with task state synchronized between the two roles in real time over WebSockets. A single pipeline carries a session from protocol definition through per-task SEQ, post-session instruments, and a semi-structured interview to an




Fig. 1. Architecture: two browser roles over a shared Postgres backend; row-level security scopes evaluator data to accounts and participant access to per-session join codes; realtime channels synchronize both views.

automatically generated PDF report. Unlike the commercial services surveyed above, Avalux is open and self-hostable (a React single-page application over Supabase Postgres with row-level security and realtime channels), and unlike prior academic tooling it makes fine-grained evaluator-side logging and evaluator-participant realtime synchronization first-class features of the moderated session rather than post hoc analyses.

III. SYSTEM DESIGN

A. Architecture

Avalux is a TypeScript React single-page application (React 19 with TanStack Router and Query, Tailwind CSS, shadcn/ui components, and Recharts visualizations, built with Vite) served statically from a CDN (Vercel), with all persistent state in a hosted Postgres instance (Supabase) accessed through PostgREST and secured by row-level security (RLS) policies (Fig. 1). Reports are rendered client-side with jsPDF, and the codebase is exercised by a unit test suite (Vitest) run on every push through continuous integration (GitHub Actions). Three backend primitives carry the design: (i) *RLS dual-role security*: evaluator rows are scoped to their authenticated account, while participants operate anonymously, authorized per session through unguessable join codes checked inside the policies themselves; (ii) *realtime channels*: the participant client subscribes to Postgres change streams over WebSockets and follows the evaluator’s mutations (task advanced, status changed) without polling, while the evaluator’s cockpit refreshes the participant-progress gate by short-interval polling of the same persisted state; (iii) *storage buckets* hold media answers captured in the participant’s browser via `getUserMedia`. There is no custom application server, which keeps the platform self-hostable on free tiers of commodity services.

B. Protocol Model

A *template* is a reusable test protocol: an ordered set of *tasks*, optionally grouped, each carrying instructions and optimal time/action baselines; a study-specific *error taxonomy* (code plus label, e.g. *WB*, *wrong button*); *per-task questions* of seven types (open text, single and multiple choice, rating, audio, video, photo); a semi-structured *interview script*; and *participant fields*: extra attributes such as handedness or vision correction that the study collects per participant. Templates also select the post-session instruments a study administers, in

any combination: the SUS [1], the unweighted (raw) NASA-TLX workload scale [49], and the eight-item UEQ-S user experience questionnaire [50]. None is mandatory, because they answer different questions. A study of a system operated under time pressure may need workload and not perceived usability; a comparison of two designs’ experiential appeal may need UEQ-S alone; a formative study run mid-iteration may need no questionnaire, its evidence being the problems observed. Opting out remains a visible decision rather than an oversight: the protocol review of Section III-J notes a template that administers no instrument at all, with the reason a standardized scale matters. A *session* instantiates a template for one participant and accumulates per-task results (status, time, action count, error and hesitation logs with timestamps, SEQ), question answers, interview answers, and the item scores of every administered instrument. Participants can be enrolled directly, through single-use personal links, or through shared links that spawn one session per respondent, with configurable collection (or omission, for anonymous participation) of identity fields. Task presentation order is a per-session strategy chosen at creation: the template’s fixed order, an independent random shuffle, or Latin-square rotation in which consecutive sessions rotate the starting task so that every task occupies every position equally often across the cohort; the strategy is recorded on the session and stated in its report.

C. Live Session and Synchronization

During a session the evaluator’s cockpit drives the protocol (Fig. 3). Completing a task opens the SEQ prompt and advances a persisted task pointer; the participant’s device follows that pointer and shows the active task’s instructions, so the facilitator never has to read them aloud. When the evaluator marks a task complete, its questions unlock on the participant’s device; conversely, the evaluator’s success/partial/failure controls are *disabled* while the participant is still answering the previous task’s questions, with a notification when they finish: a two-way gate designed to keep both roles aligned without requiring verbal coordination (Fig. 2). Both directions of the gate are computed from persisted state (the task pointer and the ratio of answers to questions per completed task), not from ephemeral messages, so either client can reload or lose connectivity mid-session and rejoin in a consistent state. The evaluator can step back one task without destroying recorded data, or explicitly reset a task (clearing its metrics, logs, and answers) after confirmation. After the last task the participant answers the interview, when the template has one, and then each instrument the template administers, on their own device; scores are computed automatically, and SUS scores are banded [27], [26].

D. Measures

For each task i the platform records completion status $s_i \in \{\text{success, partial, failure, skipped}\}$, time on task t_i , evaluator-counted actions a_i , categorized error and hesitation events with within-task timestamps, and the SEQ rating




Fig. 2. Two-way gating during a moderated session. The participant follows the evaluator over a realtime channel; the evaluator’s gate is refreshed by short-interval polling of the same persisted state, so either client can reload mid-session and rejoin consistently.

$q_i \in [1, 7]$. Templates supply per-task baselines t_i^* (optimal time) and a_i^* (optimal action count), following the optimal-path-referenced efficiency measures recommended in the measurement literature [28], [29]. Session-level indicators are the completion-status distribution, mean time and action efficiency $\frac{1}{|A|} \sum_{i \in A} t_i^*/t_i$ (and analogously for actions), where A contains only *attempted* tasks with $t_i > 0$ (skipped tasks are excluded rather than scored as zero-time successes, an aggregation rule we adopted after observing that including them produces unbounded efficiency values), and, when the SUS is administered, its score with the adjective band [1], [27]. Cross-session SUS summaries are accompanied by two-sided 95% confidence intervals computed with the t -distribution, following the small-sample recommendation of Sauro and Lewis [29]. A second rater can score a completed session’s tasks independently; the platform then reports Cohen’s κ [51] on completion status, read against the Landis and Koch bands [52], and exact-match rate, mean absolute difference, and Pearson correlation on action, error, hesitation, and SEQ counts, making the reliability of evaluator-logged measures measurable per session rather than assumed. The cockpit also records the corrections made while logging (an undone action, error, or hesitation, a return to the previous task, a task reset, a timer reset), which would otherwise leave no trace, since an undo deletes what it undoes. A reset in particular changes how a result must be read: the participant attempted the task twice, and its recorded measures describe only the second attempt, so the session view names it. The cockpit writes a marker when it opens, which lets a session logged without corrections be told apart from one that was not recorded at all. When a template administers them, raw NASA-TLX is scored as the unweighted mean of its six subscales [49] and UEQ-S as pragmatic, hedonic, and overall means on the -3 to $+3$ scale [50]. The participant client administers all instruments in English or European Portuguese, the SUS items following a validated Portuguese translation. Distributions, efficiency ratios, and SUS scores are computed at display time from the stored task results; per-task error, hesitation, and action counts are persisted when the task completes, alongside the




Fig. 3. Evaluator cockpit during a live session: task navigator with completion controls, timer, action counter, one-tap typed error logging, and hesitation logging.

individual timestamped log entries they summarize.

Because evaluator-counted actions are a manual measure, the platform can optionally instrument browser-based systems under test directly: a dependency-free script embedded in the tested system streams click, keypress, and navigation events (element descriptors and key classes only, never typed text) into the running session, keyed by its join code and accepted server-side only while the session is in progress. The evaluator cockpit shows the automatic counts beside the manual tally during the session, the results view charts the captured event stream over session time, and the events export as a raw table, enabling a manual-versus-automatic count comparison as an internal validity check on the action measure.

E. Security Model

The platform’s least conventional engineering choice is placing the whole authorization model inside the database. Evaluator tables carry policies of the form `user_id = auth.uid()`; participant access is granted to the anonymous role through policies that join the target row to a session whose unguessable `join_code` is non-null, so possession of a link is the capability to act in exactly one session. Because participants must also *write* (answers, SUS items, media references), each such table pairs a scoped `INSERT/UPDATE` policy with the read policy, and mutable invitation counters are further restricted with column-level grants so an anonymous client can increment a response counter but not rewrite an invitation’s task selection. This model survived two studies without an application server, though it requires discipline: one gap, a missing anonymous update policy on invitation counters, which silently aborted link joins after the participant row was already created, was found during operation and is now covered by a regression migration. Participant media (photo, audio, and video answers) lives in a private bucket and is served exclusively through short-lived signed URLs, so a leaked link expires instead of exposing a recording permanently; deleting a session also deletes its stored media. Studies that must avoid personal data can disable identity

collection entirely, producing anonymous participant records that never join the evaluator’s roster, and a completed session’s participant can be anonymized retroactively (identity replaced, contact details and custom attributes removed, demographics and metrics kept). Operational error monitoring follows the same stance: only uncaught exceptions are reported, tagged with opaque session identifiers, with no session replays, no performance traces, and no request payloads.

Evaluators can optionally group into organizations; a template assigned to an organization, with its sessions, becomes accessible to members through additive membership policies layered over the per-user rules, so collaboration never widens anonymous participants’ access. A restricted student role scopes selected members to individually assigned templates, supporting classroom deployments in which each student group runs its own project under a supervising account, and anyone with read access to a running session can follow it on a read-only observation view, attaching timestamped observer notes that export alongside the session data.

F. Analysis and Reporting

Session dashboards chart completion status, time and action counts against the template’s optimal baselines, error and hesitation distributions, SEQ per task, and a normalized per-task performance radar (Fig. 4); per-task event timelines plot each logged error and hesitation at its within-task timestamp, showing *when* inside a task problems clustered; template-level analytics aggregate all completed sessions. A one-click export renders a self-contained PDF report (participant demographics, per-task tables, charts, timelines, logs, interview answers, and questionnaire breakdowns with cross-session comparison), replacing the manual assembly step that typically follows moderated studies. For external statistical analysis, every measurement a study produced can also be exported as raw flat tables (CSV or JSON): sessions, task results with baselines, timestamped event logs, question answers, and per-item questionnaire scores.

G. Qualitative Coding

Beyond the quantitative record, the platform supports a lightweight qualitative analysis workflow directly on collected text. Each template carries a code book: a set of short, color-coded labels with optional definitions, authored by the evaluator. Any open-ended task answer or post-session interview answer can then be tagged with any number of codes from the session’s results view, and a code-frequency matrix (codes as rows, sessions as columns) aggregates tag counts at the template level as coding proceeds. Tags are exported as their own raw table alongside the quantitative data, with each tag row resolving the coded answer’s text, source question, and session, so downstream analysis can move between counts and quotes. Commercial platforms typically confine tagging to recorded video timelines in their unmoderated products; coupling a template-scoped code book to the structured answers of moderated sessions keeps the coding unit aligned with the protocol itself.

H. Heuristic Inspection and Synthesis

Usability testing is ordinarily preceded by inspection, in which evaluators judge an interface against a short set of principles without participants [30]. Avalux supports it as the stage before a protocol is written. A team inspects its own prototype or, before one exists, a comparable product, against a heuristic set; Nielsen’s ten heuristics [31] are built in, and the schema admits sets defined by an organization. Each finding records a description and optionally a heuristic, a location, and a severity from 0 (not a problem) to 4 (catastrophe).

Evaluators work alone. Row-level security refuses to return another evaluator’s findings until both have submitted, and submission freezes a pass so that it cannot be revised after reading the others; an organization owner who is not an evaluator sees who has submitted, but no findings, until every pass is in. The statistic that follows depends on this rule. Hertzum and Jacobsen attribute part of the evaluator effect to vague evaluation procedures leading to anchoring [32], and an evaluator who reads a teammate’s list first inflates exactly the agreement being measured. The passes are then merged into one problem list, and three figures describe the merge: any-two agreement, the mean over all pairs of evaluators of the problems they share divided by the problems they found between them, as Hertzum and Jacobsen define it; the share of problems that only one evaluator found; and the expected number of problems found by every possible group of k evaluators, computed exactly over subsets, with the number of evaluators the discovery model of Nielsen and Landauer [4] predicts for three quarters and nine tenths of the known problems. A cohort thus derives the familiar advice on how many evaluators to use from its own passes, rather than receiving it.

After testing, the same list is set against what participants did. Each predicted problem is marked as encountered, not observed, or not yet tested, and linked to the sessions that show it; problems testing revealed but inspection did not predict are recorded alongside. A session may serve as evidence only for a problem of its own study, a rule enforced in the database, and the interface does not offer pilot sessions as evidence. Two ratios follow. Of the tested predictions, the share participants encountered is the *validity* of the inspection, and of the problems testing showed, the share the inspection predicted is its *thoroughness*, measures defined by Sears [33] and systematized by Hartson et al., who make both relative to an *actual criterion* that decides which problems are real [34]. Here that criterion is the team’s own testing, and three consequences are built in. First, no prediction is ever recorded as a false alarm, because laboratory testing is not an ultimate criterion: a typical test will miss some usability problems, and its tasks are chosen by the people who designed it [34]. Second, predictions that no task exercised are left out of both ratios rather than counted against the inspection, a deliberate departure from the definition of validity, whose denominator is every problem reported. Third, the reference set is what testing showed and never the union of inspection and testing results,



Fig. 4. Session results dashboard: per-task outcomes with time, action, error, hesitation, and SEQ columns; charts below compare performance against the template's optimal baselines.

since a union would make the validity of every contributing method equal to one by construction [34].

I. Design Notes

Two design decisions proved more consequential than expected. First, *logging must not compete with observation*: every evaluator action during a task (counting an action, categorizing an error, marking a hesitation) is a single tap on a persistently visible control, sized to mobile touch-target guidelines so the cockpit is operable on a tablet held in the lap while watching the participant. Instruments that require navigation or typing (task questions, notes) are deferred to task boundaries. Second, *the protocol is an artifact, not configuration*: because a template fully encodes tasks, baselines, taxonomies, and instruments, it can be reviewed before the study like a questionnaire, reused across sessions unchanged, exported as printable observation sheets when a study must run entirely on paper, and, as in study 2, seeded from a written protocol document under version control.

J. Pedagogical Design

Avalux is also deployed in an undergraduate HCI course, where student teams each run a moderated study of a system they are designing [35], a setting in which, as Related Work notes, tools built for the purpose are rare and rarely evaluated [36]. What novices actually find hard has been characterized empirically rather than assumed: Oleson et al. identify eighteen types of learning difficulty in HCI design education, among them why an activity is done this way (WHY), how to perform a method (HOW), how to interpret feedback (SYNTH), how to avoid biasing one's own design (BIAS), when to move to the next stage (STAGE), and a tendency to rush to implementation and discount design work (RUSH) [37]. Table I names, for each objective, the difficulty it addresses, so the alignment rests on reported difficulty rather than on our own assumption about what is hard, and states the constructive alignment [53] between each objective, what the platform offers toward it, and the evidence that would show

it met; that third column is also the rubric of the classroom study in Section VI.

A tool that makes an expert faster does not thereby make a novice better; the education literature separates *productivity tools* from *mindtools*, which engage the learner in thinking about the domain [54]. Our claim is of the second kind, and the clearest way to see the difference from the traditional stopwatch and spreadsheet is in what the features cost. A productivity tool removes steps. Several of Avalux's add them: a second rater scores a session that is already scored, evaluators inspect alone before merging, each prediction is checked against the sessions, and whoever took part in a session is asked to reflect on it. Every added step produces evidence about the method rather than about the design under test, and each turns a property that a course would ordinarily assert in a lecture into a figure computed from the team's own work: that careful evaluators disagree (Cohen's κ on completion, any-two agreement across inspection passes), that one evaluator is not enough (the discovery curve of the team's own passes), that inspection and testing find different problems (the two ratios of Section III-H), and that five participants license only a wide interval (confidence intervals beside SUS). The traditional approach can state each of these, but it cannot hand a team the figure for its own study without an analysis nobody performs by hand, and an instructor who receives only a final report cannot tell whether moderation was biased, whether two observers would have logged the same errors, or whether the sample supports the conclusion. That rests on the cognitive-apprenticeship observation [55] that expert practice is invisible to novices until made visible, scaffolded, and reflected on, and the features below are grouped accordingly.

Made visible. The disagreement co-rating reveals is not a student failing. Reviewing eleven studies of cognitive walk-through, heuristic evaluation, and thinking-aloud, Hertzum and Jacobsen report that average agreement between any two evaluators of the same system ranges from 5% to 65%, and that the effect holds for experienced and novice evaluators alike and for severity ratings as well as problem detection [32]. They attribute it to vague goal analyses, vague evaluation procedures, and vague problem criteria: precisely the three things a protocol template is forced to make explicit before a session runs. Measuring κ , or the agreement of inspection passes, therefore does more than grade a pair of students; it exposes a property of the method itself.

Scaffolded. The protocol review gives feedback, checking a template without blocking against common novice mistakes (tasks that name interface labels and so lead the participant, tasks without a success criterion or optimal path, no practice task, no error taxonomy, no post-session instrument), each with a rationale linked to the tutorial. Its review appears only once a design is saved, not while it is typed: students value procedural scaffolding but distrust it when it pre-empts their own judgment [40], and a deterministic check carries none of the sycophancy they report in generative facilitators. Existing support for student studies addresses planning: a toolkit for method selection, power, and counterbalancing [38], and a

widely used plan template whose assessment turns on an instructor gate before any session may run [39]. Avalux begins where they stop, at the session itself. Its own gate, enabled per template at an organization’s discretion, is deliberately narrower than that template’s. Schwind et al. warn that close supervision calls into question the independence of a student’s research work [38], so approval gates *recruiting* and not *rehearsal*: a team may run its protocol before approval, and those sessions are recorded as pilots and excluded from the template’s aggregates. The instructor thus reviews with the protocol review’s findings and the team’s own pilot data in hand, without having authored the protocol. Every decision is kept, together with the protocol exactly as it was submitted, so two consecutive submissions are the before and after of one revision and the mistakes that survive a review can be counted rather than recalled.

Reflected on. After every completed session, each person who took part answers three fixed questions: what they did not expect, what they would change in the protocol, and what they did that may have led the participant. The questions sit beside what the session recorded (tasks failed, reset, or taking twice the expected time, hesitations, a high ease rating on a task that went badly, and notes from anyone observing), so the answer is drawn from evidence rather than from memory. A reflection stays private until submitted and cannot be revised afterwards, and a teammate’s becomes readable only once one’s own is submitted: the same rule as for inspection passes, since an account written after reading another’s is anchored to it.

For the instructor, a class view gives one row per project, with what needs attention first and, as columns, the evidence of Table I wherever the platform records it. It reads through the same row-level security as every other view, so an inspection still collecting passes contributes its progress and none of its findings. It deliberately omits the two outcomes the platform cannot observe, whether a report states its uncertainty and whether a warning was resolved with a written justification, rather than display a figure for them.

These are design arguments, not yet results. The platform acts on difficulties of method execution and interpretation, not on difficulties of design itself (WHAT, ADAPT, DIVRS) or of teamwork, scope, and stakeholders, and an added step does not guarantee that it is taken seriously. Whether the evidence these features produce changes what students understand is the question of the classroom study in Section VI.

IV. CASE STUDIES

Two studies were run end to end on the platform. Both used in-person, co-located moderated sessions in which the participant handled the system under test while the evaluator logged performance in Avalux. Study 1 used the evaluator client alone; study 2 additionally used the participant’s own device for questions, interview, and SUS via join links, so the synchronization and gating mechanism of Section III-C was exercised only in study 2’s co-located sessions; its validation with a non-co-located moderator remains future work. Participants in both studies were volunteers recruited

TABLE I
CONSTRUCTIVE ALIGNMENT: EACH LEARNING OBJECTIVE, WHAT THE PLATFORM OFFERS TOWARD IT, AND THE EVIDENCE THAT IT WAS MET. SMALL CAPITALS NAME THE LEARNING DIFFICULTY OF OLESON ET AL. [37] THAT THE OBJECTIVE ADDRESSES.

Learning objective	What Avalux offers	Evidence of outcome
Inspect before testing, with several evaluators (HOW)	Independent heuristic passes, hidden from one another until submitted	Passes submitted independently; any-two agreement reported
Say why a method step matters (WHY)	Protocol review: every check carries a rationale linked to the tutorial	Warnings resolved with a written justification
Write tasks as goals, not procedures (BIAS) Fix criteria and baselines before running (HOW)	Protocol review, leading-task check Success criterion and optimal path per task	Warnings resolved before the first session Both present on every measured task
Control order effects (HOW) Treat “error” as a judgement (SYNTH)	Order strategy per session, with explainer Co-rating, Cohen’s κ	Strategy chosen and justified in the report κ reported; disagreements reconciled
Report uncertainty with small n (SYNTH) Moderate without leading (BIAS) Analyse qualitatively as a procedure (HOW)	SUS with 95% CI Observation view, timestamped notes Code book, frequency matrix	An interval beside every score Peer notes on hints and interventions Codes defined before tagging; matrix reported
Know when a protocol is ready to run (STAGE, RUSH)	Optional instructor gate; unapproved runs recorded as pilots; every decision kept	Approval before the first recruited participant; revisions address the review
Test predictions against observation (SYNTH)	Predicted problems marked encountered or not observed, linked to sessions	Validity and thoroughness reported, untested predictions stated
Reflect on a session from its evidence (SYNTH, BIAS)	Reflection prompt beside recorded failures, resets, and observer notes	A submitted reflection for every session
Handle participant data ethically	Retroactive anonymization, private media, per-template consent recorded at join	Anonymized export; a consent record for every participant

from the authors’ university environment, were informed about the study’s purpose and the data collected (including, in study 2, browser-captured screenshots and recordings), and consented to participate. In line with the responsible institution’s guidelines, this minimal-risk usability evaluation with adult volunteers did not require formal ethics review. Participants received no compensation, gave informed consent before their session after being briefed on the study’s goals, the tasks, and the data collected (including, in study 2, screen and audio recordings), and could withdraw at any time without consequence. Collected data are stored pseudonymously, and identifying fields can be removed retroactively through the platform’s anonymization mechanism.

TABLE II
STUDY 1: SELECTED TASK-LEVEL RESULTS ($N=4$).

Task	Mean SEQ	Mean time (s)	Note
Wake up panel	1.8	—	all struggled
Recover from error	2.5	—	no error feedback
Set clock to 08:30	—	43.6	10 actions vs. 8 optimal
Read current time	7.0	2.7	100% success

TABLE III
STUDY 2: PROTOCOL STRUCTURE (TASK GROUPS AND PER-TASK QUESTION INSTRUMENTS).

Task group	Tasks	Question instruments
Study ingestion	2	choice, open, rating, photo
Data browsing	2	choice, open, rating
Areas of interest (AOI)	4	choice, rating, open, video, photo
Times of interest (TOI)	4	choice, rating, open, video, photo
Visualizations	1	choice, photo
Participant comparison	2	choice, open, rating, photo
Metrics explorer	5	choice, rating, open, audio
Export	1	choice, rating, audio

A. Study 1: Household Touch-Panel Controller

Four participants (ages 20–23, two female, three of low and one of high self-rated technical proficiency) attempted 14 tasks, seven simple (e.g. wake the panel, read the time, adjust temperature) and seven complex (e.g. set the clock, change language, recover from an error state), on a commercial water-heater touch panel. Sessions averaged 15 minutes. Aggregate completion across the 4×14 task attempts was approximately 70% success, 20% partial, and 11% failure (percentages independently rounded). SUS scores were 25.0, 40.0, 17.5, and 30.0 (mean 28.1, 95% CI [13.1, 43.1], “Poor” [27]; even the interval’s optimistic bound sits far below the benchmark mean of 68 [26]); notably the high-proficiency participant produced the *lowest* score, indicating problems not attributable to inexperience. Task-level metrics located the failures precisely (Table II): the undiscoverable wake-up gesture (SEQ 1.8, all participants struggled), error recovery without any error feedback (SEQ 2.5, participants resorted to power cycling), and clock adjustment (43.6 s and 10 actions against an optimal 8). Participants’ reports of buttons changing function across screens figured among the study’s principal design recommendations, alongside a persistently visible entry point to settings and explicit error feedback.

B. Study 2: Eye-Tracking Analysis Dashboard

The second study evaluated BEHOLD, a post-hoc analysis dashboard for Tobii Pro Glasses 3 recordings, with two participants (ages 22–23, students, low technical proficiency). The protocol exercised the platform’s richer features: 21 tasks in eight groups following the analyst’s journey (Table III), of which each session ran 20, a 19-item error taxonomy, media question types (screenshots, screen recordings, audio narration), and custom participant fields (handedness, vision correction, device, location).




Fig. 5. Per-participant SUS scores in both case studies against the industry benchmark mean of 68 [26]. Participant indices are per study; the two cohorts are unrelated.

Task success was near-universal, with one consistent exception: the fixation-metric configuration task in the metrics explorer failed for both participants, isolating the dashboard’s least intuitive workflow; every other task was completed successfully by both participants (one task was skipped in one session). Both participants scored the system at SUS 80 (“Good”, approaching “Excellent” [27]). Open answers converged on three concrete defects (an auto-track feature hidden in an overflow menu, participant IDs shown instead of names in comparisons, and truncated axis labels), each traceable to a specific task’s logs and answers in the report.

C. Cross-Study Reading

Fig. 5 contrasts the per-participant SUS scores of both studies against the industry benchmark. The two cohorts are comparable (young adults in their early twenties, predominantly low technical proficiency), yet the platform separates the two designs by more than fifty SUS points with non-overlapping ranges, and its task-level instruments explain *where* each design earns its score. We take this as evidence that an integrated moderated pipeline can deliver both the discriminative summary metrics and the diagnostic detail that typically require separate tools.

Operationally, a single evaluator ran every session in both studies. The protocol-as-template design meant the second study’s considerably richer instrumentation (19 error codes, media questions, custom participant attributes) cost setup effort once rather than per session, and both study reports, including all charts and per-task appendices, were generated without manual assembly. Running real studies also stress-tested the platform itself: study operation exposed a row-level-security gap in the anonymous invitation-counter update (Section III-E) and showed that evaluators under load can leave action counters untracked, both of which fed fixes and roadmap items back into the platform.

V. DISCUSSION AND LIMITATIONS

Construct validity

Action and error counts are logged by the evaluator rather than instrumented automatically; this keeps the platform applicable to physical devices (study 1 evaluated hardware no

browser can instrument) at the cost of observer load and potential undercounting during fast interaction bursts, and error categorization is single-observer unless a session is co-rated; the co-rating instrument measures agreement per session, but a study-level reliability figure for live logging has not yet been established. The corrections the cockpit now records give each session a direct indicator of that load, although the case-study sessions predate the feature and so have none. The synthesis ratios inherit the limits of their actual criterion: a handful of sessions, on tasks the team chose, is evidence of which predictions were encountered, not proof of which were wrong. SEQ and SUS are self-report instruments with known ceiling effects. The optimal-path baselines that anchor the efficiency measures are themselves author-defined and inherit the protocol designer’s assumptions. The protocol review can be satisfied without being understood: a student may clear its warnings mechanically, the hazard Schwind et al. note for planning toolkits [38], which is why the classroom study scores understanding and not only artifacts.

Internal validity

Both studies were run by a single evaluator who is also the platform’s author; observer bias in logging and status judgments cannot be excluded. The two case-study cohorts, while demographically similar, attempted different systems with different protocols, so the cross-study SUS contrast (Fig. 5) should be read as discriminative range, not as a controlled comparison.

External validity

The samples (four participants in study 1, two in study 2) are in line with problem-discovery practice [3], [4] but far too small for statistical generalization of the metric values (study 2’s two participants in particular make its aggregate figures illustrative only), and both cohorts were young adults in their early twenties. More importantly for a tool paper, we have not yet run a controlled comparison of evaluator effort and data quality against an ad-hoc toolchain or a commercial moderated platform; the case studies demonstrate feasibility and output quality, not superiority; that comparison is the designed next study (Section VI). Named in the vocabulary of toolkit research, the two case studies are a *demonstration*: the most common evaluation strategy in toolkit papers and the weakest on its own, since it establishes that something can be built rather than that it is worth building [41]. The evaluator and classroom studies of Section VI are *usage* studies. The two remaining strategies of that taxonomy, technical performance and a heuristic analysis of the toolkit’s own design, are not yet attempted.

Threats to the pedagogical claim

The argument of Section III-J has threats of its own, distinct from those of the case studies, and the classroom study is designed around them. The platform keeps inspection passes and reflections apart, but students who share findings in a chat are beyond its reach; the agreement of a team’s passes

detects only the gross case, since agreement far above the published range suggests the passes were not independent. The class view shows which evidence each team has produced, which invites filling a column without the understanding it stands for, so the study’s outcomes are understanding items and blind grading, never the columns themselves. Instructors read submitted reflections, which may lead students to write for the reader, so grades must not depend on what a reflection admits. Because several features add work by design, effort differs between cohorts more than it would for a tool that only saves time. A new tool also carries a novelty effect that a familiar toolchain does not. Finally, students are a captive population: taking part in the study must be separable from the assignment, and grading should involve someone independent of the tool’s development.

Security surface

The anonymous-participant model trades open link joins for RLS policies that must admit exactly the rows a code-holder is entitled to. An audit after the case studies found that the original participant-flow policies tested whether a row was the *kind* a participant may see, not whether the caller held its code, which for the anonymous role is equivalent to no policy: participant records, answers, and live invitation codes were readable without authenticating. The policies were rewritten to compare each row against the code carried in the request, verified by an automated probe, shipped with the repository, that reads every participant-facing table as an anonymous client and fails on any row returned. The gap survived two studies and a fix to a neighbouring policy; a policy set is now treated as a tested artifact, not configuration.

Because the participant client is an ordinary web page, the same protocol runs in a lab (as in both case studies), over a video call with the join link shared in chat, or asynchronously through shared links, though the evaluator-logged instruments apply only in the moderated modes. Synchronous remote moderation approaches lab quality in prior comparisons [12], [43]; verifying that the gating protocol preserves that equivalence with a remote moderator is a natural next experiment. The participant-own-device model also inherits the device diversity of field settings [45]: the mobile-first client followed directly from pilot participants joining on their phones.

VI. CONCLUSION AND FUTURE WORK

Avalux integrates protocol definition, live moderated logging, participant-side questioning on the participant’s own device, standardized instruments, analytics, and reporting into one open platform, and carried two complete usability studies whose results discriminated sharply between a poorly and a well-received design. For teaching it adds a stage before testing and a closing one after it, independent inspection and its comparison with what participants encountered, and between them features chosen for the evidence they produce about the method rather than the effort they save.

Future work proceeds on two tracks. The *evaluation* track validates the tool itself: a counterbalanced evaluator study

against a stopwatch-and-spreadsheet baseline (setup time, logging completeness, report latency, SUS and workload on the toolchain), an inter-rater reliability study of the live-logging instrument, and a manual-versus-automatic action-count comparison using the instrumentation events of Section III to bound manual logging error. The pedagogical framing of Section III-J adds a classroom study: two cohorts of one HCI course, using Avalux and the stopwatch-and-spreadsheet toolchain respectively, compared on protocol quality scored against the evidence column of Table I, inter-rater agreement between team members on a shared recorded session, the proportion of reports stating uncertainty with their estimates, and a pre/post test of methodological understanding. That test asks, among other things, whether students who saw their own evaluators' agreement understand why one evaluator is not enough, and whether those who compared predictions with testing stop reading an unobserved prediction as a false alarm; coded reflections show whether one written beside the recorded evidence names the moderator's own interventions more often than one written from memory; and the protocol review, run over each kept submission, shows which mistakes persist after feedback; it is the kind of rigorous evaluation the mapping of HCI teaching found largely absent [36]. The platform now records much of what that study needs rather than leaving it to be reconstructed: the agreement of each team's inspection passes, the ratios of predicted to encountered problems, every protocol submission with the reviewer's decision, from which the mistakes that persist across revisions can be counted, a reflection for each session, and the corrections made while logging each one, which the reliability study can set against co-rater disagreement. The remaining gaps are narrower: instructor feedback on an individual session rather than on a protocol, comparison of one protocol across design iterations, a record of the consent text each participant agreed to rather than only the moment of agreement, and evidence images on inspection findings, withheld until they can be stored under the same isolation rule as the findings themselves. The *platform* track continues with a more flexible task model; the additions described in Section III beyond the two studies' original scope were themselves motivated by those studies and shipped since they concluded.

AVAILABILITY

The platform is deployed at <https://avalux.pt>. Its source (the application, 56 SQL migrations, the seeded protocol of study 2, the tutorial, and the report generator) is available under the MIT license at <https://github.com/hugocorreia15/usertesting>. Anonymized session data from the case studies is available on request within the limits of the participants' consent.

REFERENCES

- [1] J. Brooke, "SUS: A "quick and dirty" usability scale," in *Usability Evaluation in Industry*, P. W. Jordan, B. Thomas, B. A. Weerdmeester, and I. L. McClelland, Eds. London, UK: Taylor & Francis, 1996, pp. 189–194.
- [2] J. Sauro and J. S. Dumas, "Comparison of three one-question, post-task usability questionnaires," in *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '09)*. Boston, MA, USA: ACM, 2009, pp. 1599–1608.
- [3] R. A. Virzi, "Refining the test phase of usability evaluation: How many subjects is enough?" *Human Factors*, vol. 34, no. 4, pp. 457–468, 1992.
- [4] J. Nielsen and T. K. Landauer, "A mathematical model of the finding of usability problems," in *Proceedings of the INTERACT '93 and CHI '93 Conference on Human Factors in Computing Systems (INTERCHI '93)*. Amsterdam, The Netherlands: ACM, 1993, pp. 206–213.
- [5] UserTesting, Inc., "UserTesting Human Insight Platform," 2026, commercial platform for moderated ("Live Conversation") and unmoderated video-based user research with proprietary participant panel and AI-assisted analysis; enterprise pricing undisclosed. Accessed: 2026-07-13. [Online]. Available: <https://www.usertesting.com/platform>
- [6] —, "UserZoom UX Research Platform," 2026, former UserZoom quantitative UX research platform, merged into UserTesting in 2023 under Thoma Bravo; userzoom.com permanently redirects to this page. Accessed: 2026-07-13. [Online]. Available: <https://www.usertesting.com/platform/userzoom>
- [7] Maze Design, Inc., "Maze: User Research and Testing Platform," 2026, unmoderated-first prototype and website testing with automated metrics (completion rate, misclick rate, duration, heatmaps), moderated interviews, and AI Moderator; freemium SaaS with enterprise tiers. Accessed: 2026-07-13. [Online]. Available: <https://maze.co/>
- [8] Lookback Group, Inc., "Lookback: Remote User Research Platform," 2026, moderated and unmoderated remote research with screen/camera/audio capture, live stakeholder observation rooms, timestamped notes, and AI-suggested findings; annual per-session plans from USD 299/year. Accessed: 2026-07-13. [Online]. Available: <https://www.lookback.com/>
- [9] Loop11 Pty. Ltd., "Loop11: Online Usability Testing Tool," 2026, task-based moderated and unmoderated testing of live websites and prototypes with screen/webcam/audio recording, heatmaps, clickstreams, and automatic NPS and SUS calculation; subscriptions from ca. USD 199/month. Accessed: 2026-07-13. [Online]. Available: <https://www.loop11.com/>
- [10] Useberry, "Useberry: UX Research & Usability Testing Platform," 2026, prototype-centric unmoderated testing (Figma/ProtoPie) with task completion, time-on-task, drop-off, click-tracking and user-flow analytics, plus recently added moderated interviews; freemium tiered SaaS. Accessed: 2026-07-13. [Online]. Available: <https://www.useberry.com/>
- [11] H. R. Hartson, J. C. Castillo, J. Kelso, and W. C. Neale, "Remote evaluation: The network as an extension of the usability laboratory," in *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '96)*. Vancouver, BC, Canada: ACM, 1996, pp. 228–235.
- [12] M. S. Andreasen, H. V. Nielsen, S. O. Schröder, and J. Stage, "What happened to remote usability testing? an empirical study of three methods," in *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '07)*. San Jose, CA, USA: ACM, 2007, pp. 1405–1414.
- [13] A. Bruun, P. Gull, L. Hofmeister, and J. Stage, "Let your users do the testing: A comparison of three remote asynchronous usability testing methods," in *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '09)*. Boston, MA, USA: ACM, 2009, pp. 1619–1628.
- [14] O. Alhadreti, "A comparison of synchronous and asynchronous remote usability testing methods," *International Journal of Human-Computer Interaction*, vol. 38, no. 3, pp. 289–297, 2022.
- [15] F. Paternò and G. Ballardin, "RemUSINE: A bridge between empirical and model-based evaluation when evaluators and users are distant," *Interacting with Computers*, vol. 13, no. 2, pp. 229–251, 2000.
- [16] L. Paganelli and F. Paternò, "Tools for remote usability evaluation of web applications through browser logs and task models," *Behavior Research Methods, Instruments, & Computers*, vol. 35, no. 3, pp. 369–378, 2003.
- [17] TechSmith Corporation, "Morae: Usability testing and user experience research software," Software (Recorder, Observer, Manager); discontinued, 2004, <https://assets.techsmith.com/docs/pdf-morae/SW-Usability-T-esting-with-Morae.pdf>.
- [18] J. M. C. Bastien, "Usability testing: A review of some methodological and technical aspects of the method," *International Journal of Medical Informatics*, vol. 79, no. 4, pp. e18–e23, 2010.

- [19] G. Desolda, G. Gaudino, R. Lanzilotti, S. Federici, and A. Cocco, "UTAssistant: A web platform supporting usability testing in italian public administrations," in *Proceedings of the Doctoral Consortium, Posters and Demos at CHIItaly 2017*, ser. CEUR Workshop Proceedings, vol. 1910. CEUR-WS.org, 2017, pp. 138–142.
- [20] S. Federici, M. L. Mele, R. Lanzilotti, G. Desolda, M. Bracalenti, F. Meloni, G. Gaudino, A. Cocco, and M. Amendola, "UX evaluation design of UTAssistant: A new usability testing support tool for italian public administrations," in *Human-Computer Interaction. Theories, Methods, and Human Issues (HCII 2018)*, ser. Lecture Notes in Computer Science, vol. 10901. Springer, 2018, pp. 55–67.
- [21] A. Generosi, J. Y. Villafan, L. Giraldi, S. Ceccacci, and M. Mengoni, "A test management system to support remote usability assessment of web applications," *Information*, vol. 13, no. 10, p. 505, 2022.
- [22] J. H. H. Pedersen, M. Sørensen, J. Stage, and R. T. Høegh, "Introducing asynchronous remote usability testing in practice: An action research project," in *Human-Computer Interaction – INTERACT 2021*, ser. Lecture Notes in Computer Science, vol. 12935. Springer, 2021, pp. 320–338.
- [23] International Organization for Standardization, "ISO 9241-11:2018 — ergonomics of human-system interaction — part 11: Usability: Definitions and concepts," International Organization for Standardization, Geneva, Switzerland, Standard, 2018.
- [24] A. Bangor, P. T. Kortum, and J. T. Miller, "An empirical evaluation of the system usability scale," *International Journal of Human-Computer Interaction*, vol. 24, no. 6, pp. 574–594, 2008.
- [25] J. R. Lewis, "The system usability scale: Past, present, and future," *International Journal of Human-Computer Interaction*, vol. 34, no. 7, pp. 577–590, 2018.
- [26] J. Sauro, *A Practical Guide to the System Usability Scale: Background, Benchmarks & Best Practices*. Denver, CO, USA: Measuring Usability LLC, 2011.
- [27] A. Bangor, P. Kortum, and J. Miller, "Determining what individual SUS scores mean: Adding an adjective rating scale," *Journal of Usability Studies*, vol. 4, no. 3, pp. 114–123, 2009.
- [28] B. Albert and T. Tullis, *Measuring the User Experience: Collecting, Analyzing, and Presenting UX Metrics*, 3rd ed. Cambridge, MA, USA: Morgan Kaufmann, 2022.
- [29] J. Sauro and J. R. Lewis, *Quantifying the User Experience: Practical Statistics for User Research*, 2nd ed. Cambridge, MA, USA: Morgan Kaufmann, 2016.
- [30] J. Nielsen and R. Molich, "Heuristic evaluation of user interfaces," in *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. ACM, 1990, pp. 249–256.
- [31] J. Nielsen, "Enhancing the explanatory power of usability heuristics," in *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. ACM, 1994, pp. 152–158.
- [32] M. Hertzum and N. E. Jacobsen, "The evaluator effect: A chilling fact about usability evaluation methods," *International Journal of Human-Computer Interaction*, vol. 13, no. 4, pp. 421–443, 2001, republished in vol. 15, no. 1, pp. 183–204, 2003.
- [33] A. Sears, "Heuristic walkthroughs: Finding the problems without the noise," *International Journal of Human-Computer Interaction*, vol. 9, no. 3, pp. 213–234, 1997.
- [34] H. R. Hartson, T. S. Andre, and R. C. Williges, "Criteria for evaluating usability evaluation methods," *International Journal of Human-Computer Interaction*, vol. 13, no. 4, pp. 373–410, 2001, republished in vol. 15, no. 1, pp. 145–181, 2003.
- [35] E. F. Churchill, A. Bowser, and J. Preece, "Teaching and learning human-computer interaction: Past, present, and future," *Interactions*, vol. 20, no. 2, pp. 44–53, 2013.
- [36] A. L. de Souza Lima and F. B. V. Benitti, "Let's talk about tools and approaches for teaching HCI," in *Learning and Collaboration Technologies. Designing Learning Experiences (HCII 2019)*, ser. Lecture Notes in Computer Science. Springer, 2019, vol. 11590, pp. 155–170.
- [37] A. Oleson, M. Solomon, and A. J. Ko, "Computing students' learning difficulties in HCI education," in *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*. Honolulu, HI, USA: ACM, 2020, pp. 1–14.
- [38] V. Schwind, S. Resch, and J. Sehr, "The HCI user studies toolkit: Supporting study designing and planning for undergraduates and novice researchers in human-computer interaction," in *Extended Abstracts of the 2023 CHI Conference on Human Factors in Computing Systems (CHI EA '23)*. ACM, 2023.
- [39] E. J. Rose, "Usability testing plan template: A flexible tool for planning and teaching usability," EngageCSEdu, <https://engage-csedu.org/find-resources/usability-testing-plan-template-flexible-tool-planning-and-teaching-usability>, university of Washington Tacoma. Licensed CC BY-SA. Accessed: 2026-09-11.
- [40] A. Aziz, A. Kharrufa, and I. Johnson, "Integrating GenAI into HCI pedagogy: Supporting students' learning of evaluation methods," in *Proceedings of the 8th Annual Symposium on HCI Education (EduCHI '26)*. ACM, 2026.
- [41] D. Ledo, S. Houben, J. Vermeulen, N. Marquardt, L. Oehlberg, and S. Greenberg, "Evaluation strategies for HCI toolkit research," in *Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems*. Montreal, QC, Canada: ACM, 2018.
- [42] E. McFadden, D. R. Hager, C. J. Elie, and J. M. Blackwell, "Remote usability evaluation: Overview and case studies," *International Journal of Human-Computer Interaction*, vol. 14, no. 3–4, pp. 489–502, 2002.
- [43] A. J. B. Brush, M. Ames, and J. Davis, "A comparison of synchronous remote and local usability studies for an expert interface," in *CHI '04 Extended Abstracts on Human Factors in Computing Systems*. Vienna, Austria: ACM, 2004, pp. 1179–1182.
- [44] H. B.-L. Duh, G. C. B. Tan, and V. H.-h. Chen, "Usability evaluation for mobile device: A comparison of laboratory and field tests," in *Proceedings of the 8th Conference on Human-Computer Interaction with Mobile Devices and Services (MobileHCI '06)*. Helsinki, Finland: ACM, 2006, pp. 181–186.
- [45] A. Kaikkonen, A. Kekäläinen, M. Cankar, T. Kallio, and A. Kankainen, "Usability testing of mobile applications: A comparison between laboratory and field testing," *Journal of Usability Studies*, vol. 1, no. 1, pp. 4–16, 2005.
- [46] J. T. Simon-Liedtke, W. K. Bong, T. Schulz, and K. S. Fuglerud, "Remote evaluation in universal design using video conferencing systems during the covid-19 pandemic," in *Universal Access in Human-Computer Interaction. Design Methods and User Experience (HCII 2021)*, ser. Lecture Notes in Computer Science, vol. 12768. Springer, 2021, pp. 116–135.
- [47] J. R. Hill, J. C. Brown, N. L. Campbell, and R. J. Holden, "Usability-in-place—remote usability testing methods for homebound older adults: Rapid literature review," *JMIR Formative Research*, vol. 5, no. 11, p. e26181, 2021.
- [48] M. Lobchuk, P. R. Bathi, A. Ademeyo, and A. Livingston, "Remote moderator and observer experiences and decision-making during usability testing of a web-based empathy training portal: Content analysis," *JMIR Formative Research*, vol. 6, no. 8, p. e35319, 2022.
- [49] S. G. Hart, "NASA-Task Load Index (NASA-TLX); 20 years later," in *Proceedings of the Human Factors and Ergonomics Society Annual Meeting*, vol. 50, no. 9. SAGE, 2006, pp. 904–908.
- [50] M. Schrepp, A. Hinderks, and J. Thomaschewski, "Design and evaluation of a short version of the user experience questionnaire (UEQ-S)," *International Journal of Interactive Multimedia and Artificial Intelligence*, vol. 4, no. 6, pp. 103–108, 2017.
- [51] J. Cohen, "A coefficient of agreement for nominal scales," *Educational and Psychological Measurement*, vol. 20, no. 1, pp. 37–46, 1960.
- [52] J. R. Landis and G. G. Koch, "The measurement of observer agreement for categorical data," *Biometrics*, vol. 33, no. 1, pp. 159–174, 1977.
- [53] J. Biggs, "Enhancing teaching through constructive alignment," *Higher Education*, vol. 32, no. 3, pp. 347–364, 1996.
- [54] D. H. Jonassen, *Computers as Mindtools for Schools: Engaging Critical Thinking*, 2nd ed. Upper Saddle River, NJ: Merrill/Prentice Hall, 2000.
- [55] A. Collins, J. S. Brown, and S. E. Newman, "Cognitive apprenticeship: Teaching the crafts of reading, writing, and mathematics," in *Knowing, Learning, and Instruction: Essays in Honor of Robert Glaser*, L. B. Resnick, Ed. Hillsdale, NJ: Lawrence Erlbaum Associates, 1989, pp. 453–494.